

A Kantian Analysis of Facebook's 2012 Emotional Contagion Experiment

STS Research Paper
Presented to the Faculty of the
School of Engineering and Applied Science
University of Virginia

By

Kevin Shi

May 8, 2026

On my honor as a University student, I have neither given nor received unauthorized aid on this assignment as defined by the Honor Guidelines for Thesis-Related Assignments.

ADVISOR

Benjamin J. Laugelli, Associate Professor, Department of Engineering and Society

Introduction

During a single week in January 2012, Facebook data scientists manipulated the News Feed algorithm of nearly 700,000 users to systematically filter out positive or negative emotional content with the aim of studying the viral spread of emotions across the platform (Kramer et al., 2014). This "Emotional Contagion" experiment was conducted without explicit user awareness and subsequently published in 2014, leading to a massive, global controversy regarding platform manipulation (Kramer et al., 2014). Currently, scholars and critics predominantly analyze this incident as a failure of academic research ethics, focusing heavily on how the researchers violated standard procedural protocols for informed consent (Jouhki et al., 2016). The common consensus views the harm primarily through a utilitarian or legal lens, arguing the experiment's scientific benefits did not justify the risks to unaware participants. However, this strictly procedural understanding is limited because it ignores the fundamental software architecture of the manipulation, treating the News Feed algorithm as a neutral tool rather than an active instrument of psychological control. If we continue to view the experiment solely as an issue of consent paperwork, we lose the opportunity to understand the deeper ethical implications of how predictive algorithms can alter human autonomy. I argue that the architectural deployment of the News Feed manipulation algorithm represents a deeper moral failure than the lack of informed consent because it deliberately targeted the Kantian "lower faculty" of feeling to bypass rational consideration, because corporate and academic institutions utilized bureaucratic loopholes to evade their moral duties, and because the researchers used a legal contract to prevent rational comprehension. My analysis of this case draws on Kantian duty ethics, specifically the second formulation of the categorical imperative, which dictates that one must treat humanity always as an end and never merely as a means (Kant, 1785). This framework provides the necessary

criteria to evaluate the moral actions of the software engineers independently of the experiment's scientific outcomes. To support this argument, I will analyze primary source evidence, including the original 2014 Proceedings of the National Academy of Sciences publication, Facebook's 2012 Data Use Policy, and the researchers' public statements regarding the algorithm's design.

Literature Review

Since the publication of the Facebook Emotional Contagion study, several scholars have examined the incident to evaluate its ethical and social implications. A central focus of this scholarship is how the experiment violated traditional standards of academic research and how it reduced human agency. While these analyses provide valuable insights into procedural ethics and the philosophy of technology, they are limited because they do not evaluate the specific software architecture of the News Feed algorithm as an active instrument of moral failure.

In "Facebook's Emotional Contagion Experiment as a Challenge to Research Ethics," cultural anthropologist Jouhki et al. (2016) analyze the global academic backlash to the study. They argue that the experiment highlights a critical tension between utilitarian and deontological approaches to human-subject big data research. Specifically, Jouhki et al. (2016) focus on the debate surrounding "research as manipulation" and the researchers' failure to secure proper informed consent. They conclude that current ethical frameworks, which rely heavily on procedural oversight and explicit consent agreements, are not able to handle massive-scale corporate data mining.

Similarly concerned with the impact of the experiment on users, academic philosopher Nolen Gertz (2015) examines the case through the lens of Jacques Ellul's philosophy of technology in his article "Autonomy online: Jacques Ellul and the Facebook emotional manipulation study." Gertz (2015) argues that the true problem exposed by the experiment is not

just procedural, but also existential. Because Facebook was able to manipulate hundreds of thousands of users without experiencing a decline in its active user base, Gertz asserts that the technological system's autonomy is expanding while the human users' autonomy is simultaneously diminishing (Gertz, 2015).

Although these scholars have examined the procedural ethics and the existential loss of user autonomy, they have not yet adequately considered how the deliberate design choices of the software engineers themselves violated ethical duties. Jouhki et al. (2016) treat the algorithm largely as a neutral research tool used by an organization with flawed oversight, while Gertz (2015) highlights the loss of autonomy as a broad and unavoidable consequence of modern technological systems rather than a specific, assessable moral action. To advance understanding of this case, my research addresses this gap by applying Kantian duty ethics directly to the algorithmic architecture. I will demonstrate that the experiment represents a distinct moral failure because the software's design purposefully bypassed users' rational deliberation, thereby actively subverting their autonomous moral agency.

Conceptual Framework

My analysis of the Facebook Emotional Contagion experiment draws on Kantian duty ethics, which provides a rule-based approach to evaluating moral actions independently of their ultimate consequences. This framework allows me to assess the specific moral obligations of the software engineers who designed the algorithmic manipulation, focusing on how their structural design choices respected or violated the fundamental rational capacities of the users. To establish how these moral obligations are determined, it is necessary to first outline the core principles of this ethical theory.

Duty ethics, as described by Immanuel Kant, states that human beings possess intrinsic value derived from their universal capacity for reason and rationality (Kant, 1785). Because all individuals possess this rational capacity, they must be treated with equal concern and respect. The central principle governing moral action within this framework is the categorical imperative. Specifically, the second formulation of the categorical imperative dictates that a moral actor must "act that you use humanity, whether in your own person or in the person of any other, always at the same time as an end, never merely as a means" (Kant, 1785). To treat individuals as an "end" requires respecting their autonomy and their inherent ability to engage in rational deliberation to make their own informed choices. In Kantian ethics, moral agency relies on this faculty of reason, rather than the faculties of instinct, emotion, or desire. Conversely, treating individuals simply as a "means" involves deliberately bypassing their rational faculties to use them as instruments or tools to achieve a separate, external objective. When engineers or researchers manipulate others in ways that bypass their ability to rationally engage with a situation, they undermine the fundamental moral rationality of those individuals and consequently violate their intrinsic human value (Kant, 1785).

Drawing on duty ethics, in the analysis that follows, I begin by demonstrating how the News Feed algorithm was intentionally designed to target users' emotional states rather than engage their rational faculties. I then argue that this architectural choice functionally bypassed the users' capacity for rational thinking. Finally, I show how this algorithmic bypass effectively reduced the human users to programmable components, treating them strictly as a means to achieve the corporation's data-gathering objectives and thereby violating the categorical imperative.

Analysis

Bypassing Rational Deliberation

The architectural design of the News Feed manipulation was fundamentally immoral according to Kantian duty ethics because it intentionally bypassed users' conscious awareness which ultimately denied them the opportunity to exercise rational deliberation. To understand how the platform's architecture subverted moral agency, it is necessary to examine the primary scientific report published by the engineers. In their 2014 publication in the Proceedings of the National Academy of Sciences (PNAS), the Facebook data scientists articulate the precise psychological goal of their algorithm, stating that their system facilitated "emotional contagion, leading people to experience the same emotions without their awareness" (Kramer et al., 2014).

An important consideration is the researchers' explicit inclusion of the phrase "without their awareness." This detail is significant because it demonstrates that the algorithm was not simply presenting information for users to evaluate. Instead, the software was purposefully configured to operate below the threshold of conscious human cognition. According to the second version of Kant's Categorical Imperative, humans have a moral duty to respect the rationality and autonomy of all people, treating them always as an end and never just as a means (Kant, 1785). By engineering a digital environment that altered emotional states invisibly, the developers stripped users of the requirement for Kantian autonomy, which is the conscious recognition of reality necessary to make a rational judgment.

In addition, the intentional avoidance of user rationality was an acknowledged feature of the system's deployment. In a public statement addressing the controversy, lead researcher Adam Kramer defended the methodology by noting that targeted posts "just didn't show up on some loads of Feed" (D'Onfro, 2014). This statement reveals that the software engineers fully understood, and technically relied on, the users' ignorance to achieve their experimental results.

The platform's success in proving emotional contagion relied on making the algorithm's manipulation completely hidden. This required tricking the rational mind to get a genuine, emotional reaction. By creating an algorithm that focused on emotions and deliberately bypassed conscious thinking, the engineers who made it failed in their moral duty to respect people's freedom to choose.

Counterargument

As I have argued, Facebook's deliberate architectural design to bypass user awareness fundamentally violated the Kantian duty to respect rational autonomy. Some might think, however, that algorithmic curation is simply a presentation of information. From this perspective, one could argue that filtering a News Feed does not actually subvert a user's Kantian autonomy, as the user is still fully capable of rationally deliberating on the text they do actually see and interact with. In this view, the algorithmic curation of posts is simply a presentation of information, not psychological coercion, meaning the user retains their autonomous moral agency.

But this view fails to consider the biological speed and the explicit, documented intent of the specific sorting mechanism deployed. As detailed in the official complaint filed by the Electronic Privacy Information Center (EPIC) to the Federal Trade Commission, Facebook did not just curate topics, the company intentionally altered the News Feed specifically "to elicit positive and negative emotional responses" (Electronic Privacy Information Center, 2014). The complaint also noted that the researchers "purposefully messed with people's minds" (Electronic Privacy Information Center, 2014). The algorithm was not designed to present arguments for the rational mind to evaluate, it utilized emotional contagion, a psychological process that triggers an emotional response for users that, as the researchers noted, occurs "without their awareness"

(Kramer et al., 2014). By engineering a system strictly to elicit a sub-rational emotional reaction, Facebook ensured that users could not exercise their Kantian autonomy, as the psychological manipulation occurred before rational thought could even begin.

Minimal Risk

Beyond the architectural deception of the News Feed, Facebook and its academic partners further violated the Kantian categorical imperative by relying on a utilitarian "minimal risk" defense, treating the amount and severity of harm done as the measure for morality rather than their fundamental duty to respect human autonomy. In the aftermath of the experiment, the institutions involved attempted to justify the psychological manipulation by arguing that the users were not exposed to significant danger. However, under Kantian duty ethics, the amount of harm inflicted is morally irrelevant. The moment a human being is used simply as a means to an end is when a violation occurs.

To understand how the institutions prioritized bureaucratic risk assessment over moral duty, it is important to examine the official response from the academic publishers. A few months after the study's release, the editor-in-chief of PNAS published an "Editorial Expression of Concern" to address the growing negative public reaction. In this document, the editor quotes the authors' defense regarding Cornell University's ethical oversight: "Because this experiment was conducted by Facebook, Inc. for internal purposes, the Cornell University IRB [Institutional Review Board] determined that the project did not fall under Cornell's Human Research Protection Program" (Verma, 2014). The phrase "for internal purposes" is crucial because it reveals how the university used a loophole to avoid its ethical responsibility. By categorizing the manipulation as internal corporate troubleshooting rather than human subject research, the IRB effectively classified the nearly 700,000 users as corporate data points rather than autonomous

human beings. Under Kant's imperative, institutions have a duty to protect human rationality, and by dismissing the users simply because the data was gathered "internally," the university treated the users purely as a means to generate publishable academic data.

This failure of institutional duty is further illustrated by the legal effort to locate any entity willing to take moral responsibility for the users. In July 2014, legal scholar James Grimmelman issued a formal letter to Facebook's General Counsel, requesting the release of Facebook's own internal IRB minutes under Maryland law (Grimmelmann & Henry, 2014). Grimmelman explicitly highlighted the circular nature of this ethical oversight. He pointed out that Cornell avoided responsibility by characterizing the experiment as "research ... conducted independently by Facebook," allowing them to argue that "it was Facebook's responsibility to provide IRB review" (Grimmelmann & Henry, 2014). This exchange shows that both institutions focused on legal definitions rather than ethical duties. Under Kantian duty ethics, institutions have a requirement to respect the autonomy of human subjects, and instead of fulfilling this duty, Cornell and Facebook simply shifted the legal liability to each other. By doing this, they treated the users as a means to protect their own institutions from accountability, rather than as autonomous individuals. A strict Kantian framework demands that every individual be treated as an end in themselves. However, this evidence demonstrates that the users were treated as objects of legal liability, passed between institutions that refused to acknowledge their moral autonomy.

This risk-shifting behavior relies on a fundamentally flawed application of federal research standards. The United States standard for human subject protection, outlined in 45 CFR 46, allows for expedited review or altered consent procedures for research involving "no more than minimal risk" (Protection of human subjects, 45 C.F.R. § 46). While Facebook and Cornell relied on this utilitarian concept of minimal risk to justify bypassing informed consent, this

defense is contrary to duty ethics. Kantian ethics rejects the idea of a sliding scale of harm.

Whether the emotional manipulation was temporary, minor, or "minimal risk," the developers intentionally altered the psychological states of hundreds of thousands of people without granting them the rational choice to participate. By prioritizing a utilitarian assessment of minimal harm over the fundamental duty to obtain rational consent, the researchers failed to uphold the ethical standards required of their profession.

Consent and the Terms of Service

Facebook's reliance on its Data Use Policy as a substitute for explicit permission was a moral failure under duty ethics because the legal document was architecturally designed to prevent user comprehension, treating individuals as legal shields rather than autonomous decision-makers. When the public backlash regarding the emotional contagion experiment began, the researchers defended their actions by pointing to the platform's Terms of Service. In their official publication, the authors explicitly stated that the experiment "was consistent with Facebook's Data Use Policy, to which all users agree prior to creating an account on Facebook, constituting informed consent for this research" (Kramer et al., 2014). By treating the click of an "I Agree" button as "informed consent," the researchers attempted to fulfill their moral requirement to respect user autonomy. However, under Kantian duty ethics, a true moral contract requires that both parties possess the ability to rationally deliberate and freely agree to the conditions.

Evaluating Facebook's Terms of Service against the ethical standards for human research reveals why it fails the Kantian standard for rational deliberation. *The Belmont Report*, the primary ethical guideline for human subject research in the United States, explicitly divides informed consent into three necessary components, which are "information, comprehension, and

voluntariness" (The National Commission for the Protection of Human Subjects of Biomedical and Behavioral Research, 1979). While Facebook claimed they provided "information" through their Data Use Policy, the company actually failed to even include the words "research" or "troubleshooting" in the document until four months after the experiment concluded (Electronic Privacy Information Center, 2014). Furthermore, by structurally designing the platform's onboarding process around thousands of words of dense, non-negotiable legal jargon, Facebook actively prevented the second requirement, "comprehension" (Facebook, 2012). Presenting users with thousands of words of dense, non-negotiable legal jargon does not encourage rational comprehension. It instead relies on the unlikely scenario of the user actually reading the text. By leveraging a legal contract that users are expected to blindly accept, Facebook denied users the comprehension required to make an autonomous, rational choice about participating in psychological research.

Furthermore, documentary evidence proves that Facebook was already legally aware that its standard Terms of Service were deceptive and insufficient for acquiring true consent. In 2011, just months before the emotional contagion experiment was conducted in January 2012, Facebook signed a formal settlement with the United States Federal Trade Commission (Federal Trade Commission, 2011). The Federal Trade Commission had investigated Facebook for deceptive privacy practices, and the resulting legal agreement mandated that the corporation stop misleading users and acquire "affirmative express consent" before sharing certain data (Federal Trade Commission, 2011). The FTC's specific use of the phrase "affirmative express consent" is a corrective measure against Facebook's standard practices. This legal mandate demonstrates that the US government had already recognized that Facebook's overarching Data Use Policy did not constitute a valid, rationally understood agreement.

In spite of this 2011 mandate, the researchers proceeded to manipulate the News Feeds of nearly 700,000 users in 2012, retroactively using the same criticized Terms of Service as their ethical shield. Under the Categorical Imperative, one must "treat himself and all others never merely as means but always at the same time as ends in themselves" (Kant, 1785). By hiding their psychological experiment within a legal document that they knew users would not comprehend, which additionally the FTC had already deemed deceptive in other contexts, the researchers misused the Terms of Service. Facebook used the document primarily to legally shield the corporation, rather than to inform users. This treatment of users as a legal means to an end undermines the conscious, rational autonomy required by duty ethics.

Conclusion

The architectural deployment of Facebook's 2012 News Feed manipulation algorithm represents a fundamental moral failure under Kantian duty ethics, moving beyond a breach of informed consent. This experiment systematically undermined users' capacity for autonomous moral agency in three distinct ways. First, the algorithm's architecture was intentionally designed to operate below the threshold of conscious awareness, bypassing the rational deliberation required for human autonomy. Second, corporate and academic institutions relied on a flawed utilitarian "minimal risk" defense and bureaucratic loopholes to waive their fundamental moral duty to their human subjects. Finally, by burying the experiment within an almost certainly unread Terms of Service agreement, the researchers leveraged a legal contract to protect the corporation while actively preventing the rational user comprehension necessary for true consent.

The significance of this argument lies in how it redefines the ethical evaluation of digital platforms. By shifting the focus away from procedural checklists and utilitarian risk assessments, readers gain a new understanding of software architecture as a potential instrument of

psychological coercion. Recognizing that algorithms can be designed to exploit sub-rational biological responses means society should view digital manipulation as a challenge to human rationality, rather than only a privacy violation. Consequently, this understanding demands a shift in professional engineering practice. Software developers and data scientists must recognize that legal disclaimers and "industry-standard" testing do not substitute for moral duty. Ethical engineering in the digital age requires designing systems that structurally preserve, rather than systematically dismantle, the user's autonomous rational agency.

Word Count: 3232

References

- D'Onfro, J. (2014, June 30). *Adam Kramer Facebook mood manipulation*. Business Insider.
<https://www.businessinsider.com/adam-kramer-facebook-mood-manipulation-2014-6>
- Electronic Privacy Information Center. (2014). *In the matter of Facebook, Inc.: Complaint, request for investigation, injunction, and other relief* [Complaint to the Federal Trade Commission].
<https://epic.org/wp-content/uploads/privacy/ftc/facebook/Facebook-Study-Complaint.pdf>
- Facebook. (2012, January 4). *Facebook data use policy*.
https://web.archive.org/web/20120104214758/http://www.facebook.com/full_data_use_policy
- Federal Trade Commission. (2011). *In the matter of Facebook, Inc.: Complaint and agreement containing consent order* [Agreement containing consent order].
<https://www.ftc.gov/sites/default/files/documents/cases/2011/11/111129facebookagree.pdf>
- Gertz, N. (2015). Autonomy online: Jacques Ellul and the Facebook emotional manipulation study. *Research Ethics*, 12(1), 55–61. <https://doi.org/10.1177/1747016115579534>
- Grimmelmann, J., & Henry, L. M. (2014). *Letter to Facebook, Inc., regarding institutional review board minutes for Facebook's human subjects research* [Letter to Facebook, Inc.]. University of Maryland Francis King Carey School of Law.
<https://james.grimmelmann.net/files/legal/facebook/Facebook.pdf>
- Jouhki, J., Lauk, E., Penttinen, M., Sormanen, N., & Uskali, T. (2016). Facebook's emotional contagion experiment as a challenge to research ethics. *Media and Communication*, 4(4), 75. <https://doi.org/10.17645/mac.v4i4.579>

Kant, I. (1785). *Groundwork of the metaphysics of morals*. Cambridge University Press.

<https://cpb-us-w2.wpmucdn.com/blog.nus.edu.sg/dist/c/1868/files/2012/12/Kant-Groundwork-ng0pby.pdf>

Kramer, A. D. I., Guillory, J. E., & Hancock, J. T. (2014). Experimental evidence of massive-scale emotional contagion through social networks. *Proceedings of the National Academy of Sciences*, *111*(24), 8788–8790. <https://doi.org/10.1073/pnas.1320040111>

Protection of human subjects, 45 C.F.R. § 46 (2024).

<https://www.ecfr.gov/current/title-45/subtitle-A/subchapter-A/part-46>

The National Commission for the Protection of Human Subjects of Biomedical and Behavioral Research. (1979, April 18). *The Belmont Report*. U.S. Department of Health and Human Services.

<https://www.hhs.gov/ohrp/regulations-and-policy/belmont-report/read-the-belmont-report/index.html>

Verma, I. M. (2014). Editorial expression of concern: Experimental evidence of massivescale emotional contagion through social networks. *Proceedings of the National Academy of Sciences*, *111*(29), 10779–10779. <https://doi.org/10.1073/pnas.1412469111>